

# Task 6: K-Nearest Neighbors (KNN) Classification

## Objective

Understand and implement KNN for classification problems.

---

## Steps & Implementation

### 1. Dataset

- Dataset Used: **Iris Dataset** (scikit-learn / Kaggle)
- Target: Species (Setosa, Versicolor, Virginica)

### 2. Data Preprocessing

- Features normalized using `StandardScaler`
- Train-test split (80:20)

### 3. KNN Classifier

- Implemented using `KNeighborsClassifier`
- Tried different values of K (3, 5, 7, 9)

### 4. Model Evaluation

- Metrics: Accuracy, Confusion Matrix
- Example results (K=5):
- Accuracy: ~97%
- Confusion Matrix:

```
[[10  0  0]
 [ 0  9  1]
 [ 0  0 10]]
```

### 5. Decision Boundaries

- Visualized using matplotlib (2D feature pairs)
  - Showed how classification regions change with different K values
- 

## Interview Questions & Answers

1. How does the KNN algorithm work?
2. Classifies a data point based on majority vote of its K nearest neighbors using a distance metric.
3. How do you choose the right K?

4. Small  $K \rightarrow$  high variance (overfitting), Large  $K \rightarrow$  high bias (underfitting). Usually odd  $K$  to avoid ties.


**5. Why is normalization important in KNN?**


6. Distance metrics are scale-sensitive; normalization ensures fair contribution of all features.

**7. What is the time complexity of KNN?**

8. Training:  $O(1)$ , Prediction:  $O(n \times d)$  per query ( $n$  = samples,  $d$  = dimensions).

**9. What are pros and cons of KNN?**

10.  Pros: Simple, no training, works well on small datasets.

11.  Cons: Slow on large datasets, sensitive to noise & irrelevant features.

**12. Is KNN sensitive to noise?**

13. Yes, mislabeled or noisy data can mislead neighbor voting.

**14. How does KNN handle multi-class problems?**

15. Majority voting among neighbors naturally extends to multiple classes.

**16. What's the role of distance metrics in KNN?**

17. Defines neighbor closeness (e.g., Euclidean, Manhattan, Minkowski).

---

## Code (Colab / Jupyter Notebook)

```
# KNN Classification - Iris Dataset
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
import matplotlib.pyplot as plt
import seaborn as sns

# 1. Load dataset
data = load_iris()
X = pd.DataFrame(data.data, columns=data.feature_names)
y = pd.Series(data.target)
```

```

# 2. Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42, stratify=y)

# 3. Normalize features
scaler = StandardScaler()
X_train_s = scaler.fit_transform(X_train)
X_test_s = scaler.transform(X_test)

# 4. Train KNN with different K values
for k in [3, 5, 7, 9]:
    clf = KNeighborsClassifier(n_neighbors=k)
    clf.fit(X_train_s, y_train)
    y_pred = clf.predict(X_test_s)
    acc = accuracy_score(y_test, y_pred)
    print(f'K={k} -> Accuracy: {acc:.3f}')
    print(confusion_matrix(y_test, y_pred))

# 5. Final model (K=5)
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_s, y_train)
y_pred = knn.predict(X_test_s)

print("\nClassification Report:\n", classification_report(y_test, y_pred))

# 6. Confusion Matrix Heatmap
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('KNN Confusion Matrix (K=5)')
plt.show()

# 7. Decision Boundaries (2 features for visualization)
from matplotlib.colors import ListedColormap

X_vis = X.iloc[:, :2].values
y_vis = y.values
X_train_v, X_test_v, y_train_v, y_test_v = train_test_split(X_vis, y_vis,
test_size=0.2, random_state=42, stratify=y_vis)
scaler_v = StandardScaler()
X_train_v = scaler_v.fit_transform(X_train_v)
X_test_v = scaler_v.transform(X_test_v)

clf_vis = KNeighborsClassifier(n_neighbors=5)
clf_vis.fit(X_train_v, y_train_v)

# Mesh grid
gx, gy = np.meshgrid(np.arange(X_train_v[:,0].min()-1, X_train_v[:,0].max()
+1, 0.01),

```

```

np.arange(X_train_v[:,1].min()-1, X_train_v[:,1].max()
+1, 0.01))

Z = clf_vis.predict(np.c_[gx.ravel(), gy.ravel()])
Z = Z.reshape(gx.shape)

plt.contourf(gx, gy, Z, alpha=0.3,
cmap=ListedColormap(['red', 'green', 'blue']))
plt.scatter(X_train_v[:,0], X_train_v[:,1], c=y_train_v, edgecolor='k',
marker='o')
plt.title('KNN Decision Boundary (K=5, 2 features)')
plt.show()

```

## Repository Structure

📁	KNN-Classification-Task	
├	📄 knn_classification.ipynb	# Code
├	📄 README.md	# Explanation
├	📄 outputs.png	# Confusion Matrix, Decision Boundary
└	📁 data	# Dataset

## Key Learning

- KNN is an **instance-based learning algorithm** (no explicit training).
- Choice of K and feature scaling significantly affect performance.
- KNN works well for small, low-dimensional datasets but struggles with large-scale/high-dimensional data.